

[POC] Architecture and Design of the Auto-CAM-Prototype1: Data Flow and Key Decision

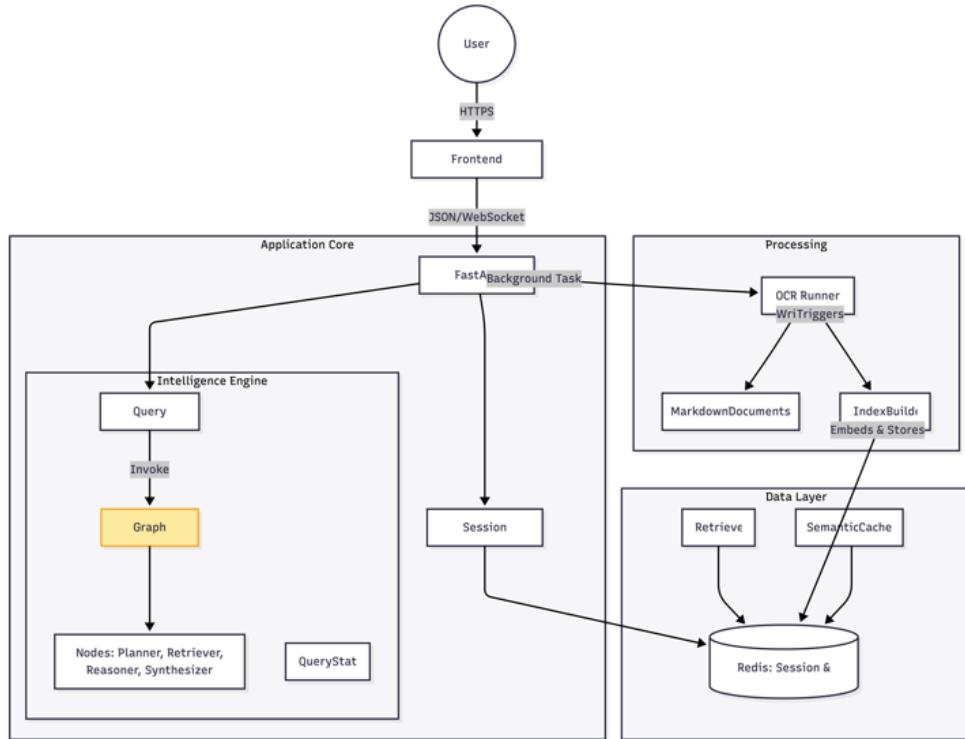


Figure 1: High Level Design

🚀 Detailed Data Flow

1. Ingestion & Digitization (The "Eyes")

The system begins by ingesting raw financial documents (PDFs, Excel).

- **Component:** **OCR Runner**
- **Process:**
 - a. **Detection:** Scans directory for files.
 - b. **Digitization:** Uses specialized engines (Google DocAI for PDFs, Custom Parsers for Excel) to extract text.

- c. **Normalization:** Converts everything into clean **Markdown** format with structure preservation.
- d. **Output:** `processed_output/*.md` files.

2. Indexing & Storage (The "Library")

We organize text so the AI can find it instantly.

- **Components:**

IndexBuilder,

Chunker, **Embedder**,

Redis

- **Process:**

- a. **Chunking:**

MarkdownChunker splits documents into small, logical "chunks" (paragraphs, tables) that fit in the AI's context window.

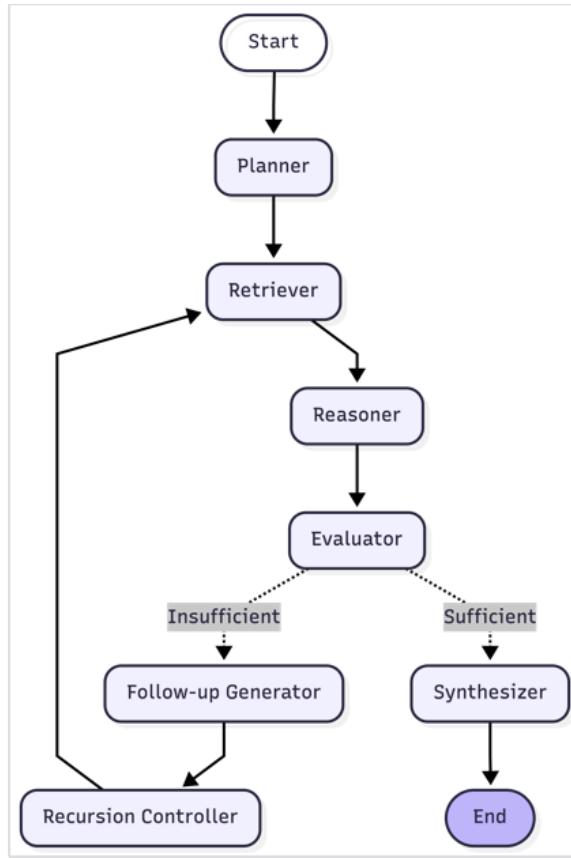
b. **Embedding:** Chunks are converted into vectors (lists of numbers) using **TITAN Embeddings v2**.

c. **Compression & Storage:**

- **Vectors:** We compress vectors using **Base64 Encoded Numpy Bytes** before storing in Redis.
- **Technique:** `base64.b64encode(vec_np.tobytes())`.
- **Advantage:** This is **3x-4x smaller** and faster to transmit than storing vectors as plain text lists (JSON).

3. The Intelligence Engine (The "Brain")

This is the core agentic workflow powered by **LangGraph**.



- **Logic Flow:**

- Planner:** Breaks the user's complex question into specific sub-steps.
- Retriever:** Searches Redis for relevant chunks (Vector Search).
- Reasoner:** Analyzes chunks to answer specific steps.
- Evaluator:** "Did we answer the question?" If no, it triggers **Recursion** (Loop back).
- Synthesizer:** Compiles the final answer.

🧠 Key Design Decisions

Why Redis? (Session Isolation & Speed)

We use Redis as our primary data layer instead of a traditional SQL database.

Advantages:

1. Session Isolation (Privacy High):

- Every key is prefixed with `session_id`.
- **Result:** Complete data separation. Session A cannot accidentally touch Session B's data.
- **Security:** Wiping a session is as simple as deleting all keys with that prefix.

2. No DB Connections:

- Redis is in-memory and extremely fast.

- We don't manage complex connection pools or schema migrations.
- This makes the system **stateless** and easy to scale (just add more Redis nodes).

Disadvantages:

- **Cold Start:** Every new session starts from scratch. We must re-index documents for every new "Session ID". (Trade-off for perfect privacy).

High Quality via Recursive Retrieval

Our system delivers exceptionally high-quality answers even from heterogeneous data (long PDFs mixed with Excel).

- **The Secret: Recursive Feedback Loops.**
- **How it works:**
 - a. The **Planner** breaks a vague question ("Assess the borrower") into specific queries ("Find Father's Name", "Check Credit Score").
 - b. The **Evaluator** checks the answer. If the retrieval missed something (e.g., hidden in a footnote), it forces the agent to **try again** with a better search query.
- **Trade-off: Time.** This "thinking" process takes longer (multiple LLM calls) but ensures accuracy.

Solving the "Long Context" Problem

LLMs have a limit on how much text they can read (Context Window). Financial documents often exceed this.

- **Solution: External Memory (Vector Retrieval).**
- **Mechanism:** instead of feeding the *entire* 100-page loan agreement to the LLM, we use Redis to find only the **top 5 relevant paragraphs** and feed those.
- **Result:** We can chat with documents of **infinite length** because we only load what is needed for the specific question.